

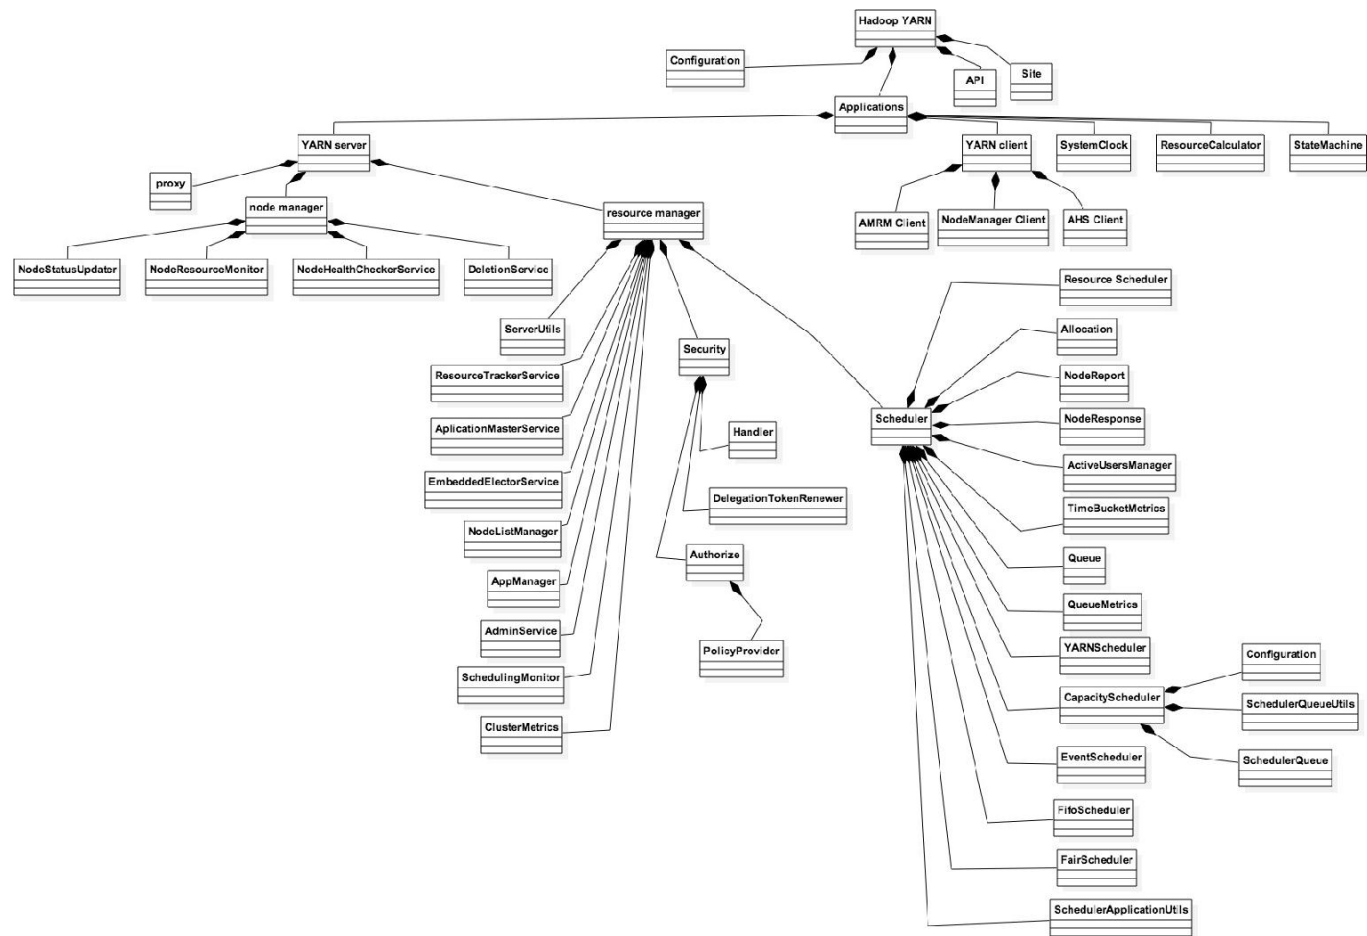
**Final project**  
**From source code to design models**  
**Hadoop Yarn Architecture**  
**Software Modeling SWEN 705-01**

**Presented by:**

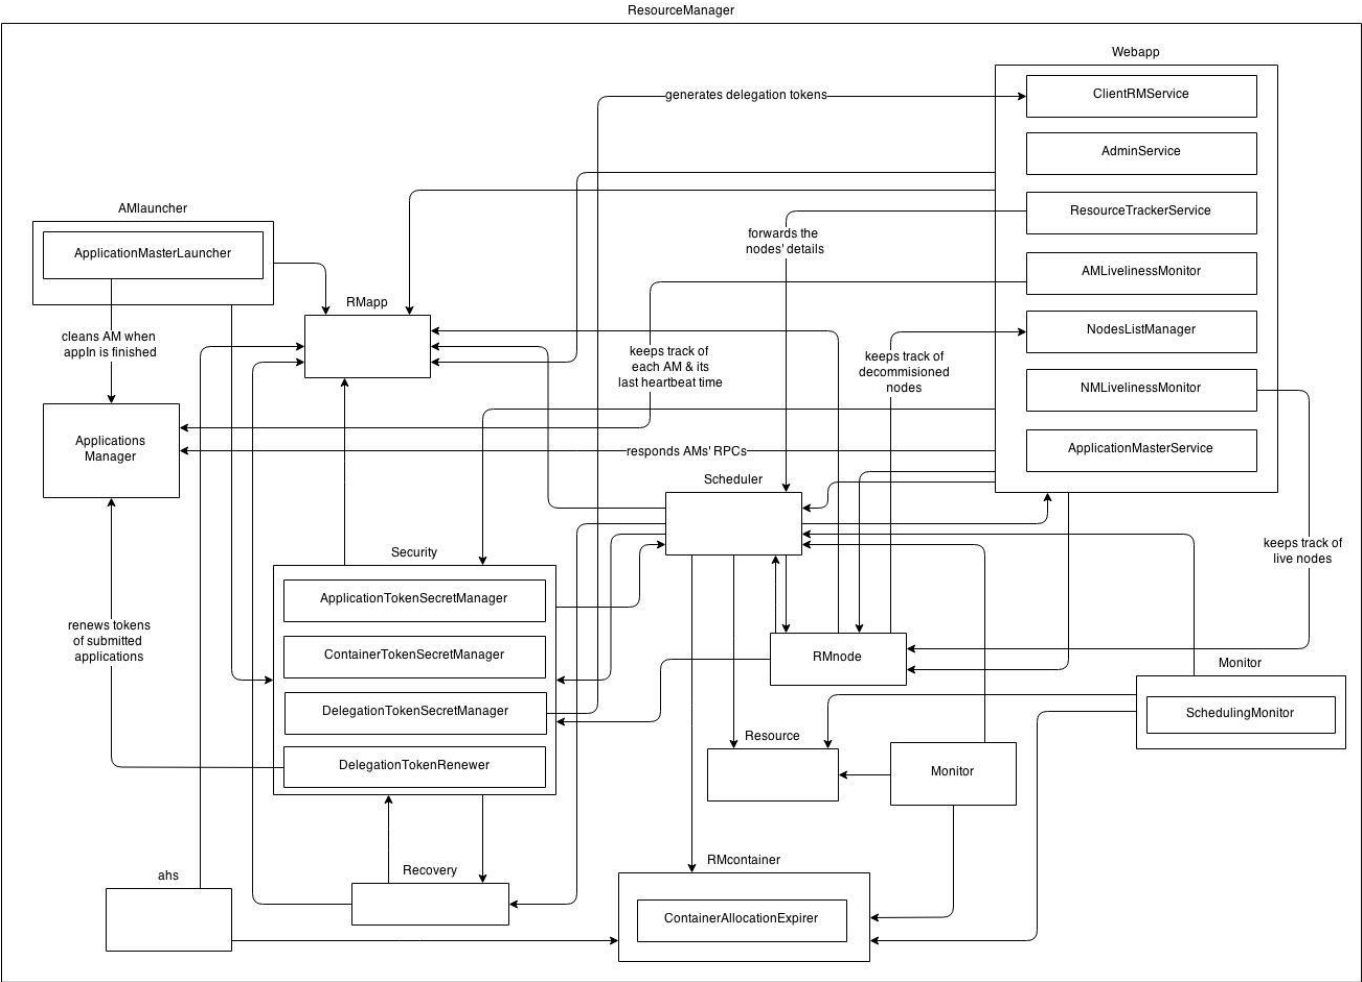
Ivan Tactuk  
Omeya Jadhav  
Yue Hua  
Leonardo Matos

Module views of the system

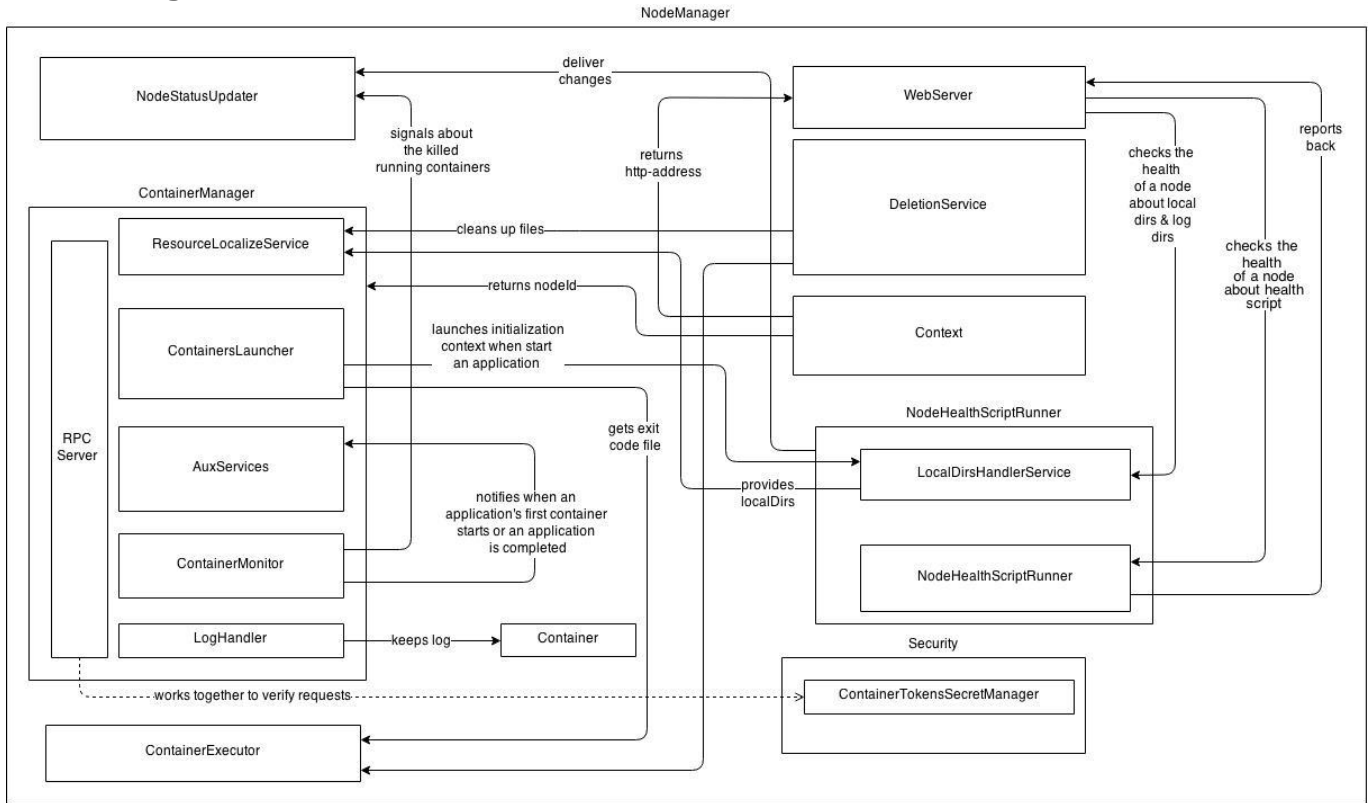
Main system model view



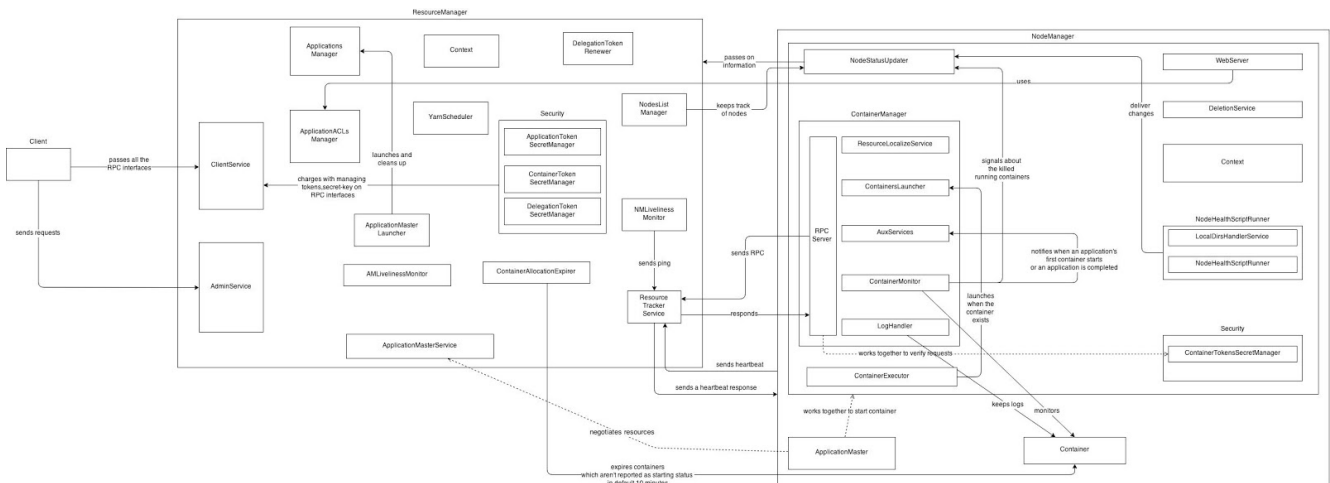
# Resource Manager module view



## Node Manager module view



## Resource Manager and Node Manager module interaction view



## Main Modules table

| Component name           | Belongs to package | Location of file   |
|--------------------------|--------------------|--|
| configuration            | Hadoop yarn        | /hadoop-yarn/conf  |
| api                      | Hadoop yarn        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-api   |
| site                     | Hadoop yarn        | Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-site   |
| applications             | Hadoop yarn        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-applications  |
| Yarn server              | applicatio ns      | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server  |
| proxy                    | Yarn server        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server  |
| Node Manager             | Yarn server        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager   |
| NodeStatusUpdater        | Webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/main/java/org/apache/hadoop/yarn/server/nodemanager/NodeStatusUpdater.java                |
| NodeResourceMonitor      | Webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/main/java/org/apache/hadoop/yarn/server/nodemanager/NodeResourceMonitor.java              |
| NodeHealthCheckerService | Webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/main/java/org/apache/hadoop/yarn/server/nodemanager/NodeHealthCheckerService.java         |
| DeletionService          | Webapp             | /hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/main/java/org/apache/hadoop/yarn/server/nodemanager/DeletionService.java                                    |
| ResourceManager          | Yarn server        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager                               |
| ServerUtils              | Webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/RMServerUtils.java            |
| ResourceTrackerService   | Webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/ResourceTrackerService.java   |
| ApplicationMasterService | webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/ApplicationMasterService.java |
| EmbeddedElectorService   | webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/EmbeddedElectorService.java   |
| NodeListManager          | webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/NodesListManager.java         |
| AppManager               | webapp             | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/RMAppManager.java             |

|                        |                  |   |
|------------------------|------------------|---|
| AdminService           | webapp           | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/AdminService.java                        |
| SchedulingMonitor      | monitor          | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/monitor/SchedulingMonitor.java           |
| ClusterMetrics         | webapp           | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/ClusterMetrics.java                      |
| Security               | resource manager | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/test/java/org/apache/hadoop/yarn/server/resource-manager/security                                 |
| Handler                | Security         | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/security/RMAuthenticationHandler.java    |
| DelegationTokenRenewer | Security         | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/security/DelegationTokenRenewer.java     |
| Authorize              | Security         | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/security/authorize                       |
| Policyprovider         | Authorize        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/security/authorize/RMPolicyProvider.java |
| Scheduler              | resource manager | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/scheduler                                |
| ResourceScheduler      | Scheduler        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/scheduler/ResourceScheduler.java         |
| Allocation             | Scheduler        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/scheduler/Allocation.java                |
| NodeReport             | Scheduler        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/scheduler/NodeReport.java                |
| NodeResponse           | Scheduler        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/scheduler/NodeResponse.java              |
| ActiveUserManager      | Scheduler        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/scheduler/ActiveUsersManager.java        |
| TimeBucketMetrics      | Scheduler        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/scheduler/TimeBucketMetrics.java         |
| Queue                  | Scheduler        | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resource-manager/src/main/java/org/apache/hadoop/yarn/server/resource-manager/scheduler/Queue.java                     |

|                           |           |   |
|---------------------------|-----------|---|
| QueueMetrics              | Scheduler | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/QueueMetrics.java                            |
| YarnScheduler             | Scheduler | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/YarnScheduler.java                           |
| CapacityScheduler         | Scheduler | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacityScheduler.java              |
| SchedulerQueueUtils       | Capacity  | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java                   |
| Configuration             | Capacity  | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java |
| Queue                     | Capacity  | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueue.java                        |
| EventScheduler            | Event     | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/event/SchedulerEvent.java                    |
| SchedulerApplicationutils | fifo      | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/SchedulerAppUtils.java                       |
| FifoScheduler             | fifo      | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/fifo/FifoScheduler.java                      |
| FairScheduler             | Fair      | /Hadoop_SourceCode/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/fair/FairScheduler.java                      |

## II) OCL and test cases:

### 1. The next heart beat interval value should be a valid positive number.

**Context** ResourceTrackerService::serviceInit()

**Inv:** nextHeartBeatInterval > 0

#### Test Case 1.1:

```
nextHeartBeatInterval=-1
try
{
    ResourceTrackerService.serviceInit()
    Assert.Fail()
}
catch (YarnRuntimeException)
{
}
}
```

**Test Case 1.2:**

```

nextHeartBeatInterval=0
try
{
    ResourceTrackerService.serviceInit()
    Assert.Fail()
}
catch (YarnRuntimeException)
{
}

```

**Test Case 1.3:**

```

nextHeartBeatInterval=2
try
{
    ResourceTrackerService.serviceInit()
}
catch (YarnRuntimeException)
{
    Assert.Fail()
}

```

**1. The Minimum allocation memory should be lower or equal than the maximum****Context** FairScheduler::validateConf()**Inv:** RM\_SCHEDULER\_MINIMUM\_ALLOCATION\_MEMORY <= RM\_SCHEDULER\_MAX\_ALLOCATION\_MEMORY**Test Case 2.1:**

```

RM_SCHEDULER_MINIMUM_ALLOCATION_MEMORY=2
RM_SCHEDULER_MAX_ALLOCATION_MEMORY=55
try
{
    FairScheduler.validateConf()
}
catch (YarnRuntimeException)
{
    Assert.Fail()
}

```

**Test Case 2.2:**

```

RM_SCHEDULER_MINIMUM_ALLOCATION_MEMORY =55
RM_SCHEDULER_MAX_ALLOCATION_MEMORY=22
try
{
    FairScheduler.validateConf()
    Assert.Fail()
}

```



```
catch(YarnRuntimeException)
{
}
```

#### Test Case 2.3:

```
RM_SCHEDULER_MINIMUM_ALLOCATION_MEMORY =55
RM_SCHEDULER_MAX_ALLOCATION_MEMORY=55
try
{
    FairScheduler.validateConf()
}
catch(YarnRuntimeException)
{
    Assert.Fail()
}
```

#### 1. The Resource scheduler minimum allocation memory should be a whole number value.

**Context** FairScheduler::validateConf()

**Inv:** RM\_SCHEDULER\_MINIMUM\_ALLOCATION\_MEMORY >= 0

#### Test Case 3.1:

```
RM_SCHEDULER_MINIMUM_ALLOCATION_MEMORY =-1
try
{
    FairScheduler.validateConf()
    Assert.Fail()
}
catch(YarnRuntimeException)
{
}
```

#### Test Case 3.2:

```
RM_SCHEDULER_MINIMUM_ALLOCATION_MEMORY =0
try
{
    FairScheduler.validateConf()
}
catch(YarnRuntimeException)
{
    Assert.Fail()
}
```

#### Test Case 3.3:

```
RM_SCHEDULER_MINIMUM_ALLOCATION_MEMORY =2
try
{
    FairScheduler.validateConf()
}
```

```

catch (YarnRuntimeException)
{
    Assert.Fail()
}

```

1. **The expire interval of the resource Manager should be equal or larger than the heart beat interval.**

**Context** ResourceManager::validateConfigs()  
**Inv:** expireIntvl >= heartbeatIntvl

**Test Case 4.1:**

```

heartbeatIntvl = 2
expireIntvl = 55
try
{
    ResourceManager.validateConfigs()
}
catch (YarnRuntimeException)
{
    Assert.Fail()
}

```

**Test Case 4.2:**

```

heartbeatIntvl=55
expireIntvl =22
try
{
    ResourceManager.validateConfigs()
    Assert.Fail()
}
catch (YarnRuntimeException)
{
}

```

**Test Case 4.3:**

```

heartbeatIntvl=55
expireIntvl =55
try
{
    ResourceManager.validateConfigs()
}
catch (YarnRuntimeException)
{
    Assert.Fail()
}

```

1. **The global Max attempts of the resource manager should be a positive number.**

**Context** ResourceManager::validateConfigs()

**Inv:** globalMaxAppAttempts>0

**Test Case 5.1:**

```
globalMaxAppAttempts=-1
try
{
    ResourceManager.validateConfigs()
    Assert.Fail()
}
catch (YarnRuntimeException)
{
}
```

**Test Case 5.2:**

```
globalMaxAppAttempts=0
try
{
    ResourceManager.validateConfigs()
    Assert.Fail()
}
catch (YarnRuntimeException)
{
}
```

**Test Case 5.3:**

```
globalMaxAppAttempts=2
try
{
    ResourceManager.validateConfigs()
}
catch (YarnRuntimeException)
{
    Assert.Fail()
}
```

**1. A valid container token should be used for correct node.**

**Context** NMContainerTokenSecretManager::retrievePassword()

**Inv:** identifier.getNmHostAddress() == nodeHostAddr

**Test Case 6.1:**

```
identifier.setNmHostAddress("197.0.0.1");
nodeHostAddr = "197.0.0.1"
try
{
    NMContainerTokenSecretManager.retrievePassword()
}
catch (SecretManager.InvalidToken)
{
}
```

```

        Assert.Fail()
    }

Test Case 6.2:
identifier.setNmHostAddress("197.0.0.1");
nodeHostAddr = "192.168.1.1"
try
{
    NMContainerTokenSecretManager.retrievePassword()
    Assert.Fail()
}
catch (SecretManager.InvalidToken)
{
}

```

# **1. When moving an Application to a Queue leaf, NumRunnableApps cannot be greater than maxRunningApps or maxShare**

**Context** FairScheduler::moveApplication(ApplicationId appld, String queueName)

**Inv:** queueMgr.getLeafQueue(queueName).getNumRunnableApps() < maxRunningApps

## **Test Case 7.1:**

```

maxRunningApps=3
try
{
    moveApplication(new Application(id=1), queueName="RIT")
    moveApplication(new Application(id=2), queueName="RIT")
    moveApplication(new Application(id=3), queueName="RIT")
    moveApplication(new Application(id=4), queueName="RIT")
    Assert.Fail()
}
catch (YarnException)
{
}

```

## **Test Case 7.2:**

```

maxRunningApps=3
try
{
    moveApplication(new Application(id=1), queueName="RIT")
    moveApplication(new Application(id=2), queueName="RIT")
    moveApplication(new Application(id=3), queueName="RIT")
}
catch (YarnException)
{
    Assert.Fail()
}

```

# **1. CapacityScheduler cannot have two leaf queues with the same name.**

**Context** CapacityScheduler  
**Inv:** Leafqueues.isUnique(Name)

**Test Case 8.1:**

```
Leafqueues.add(name="A1")
Leafqueues.add(name="A1")
try
{
    CapacityScheduler.parseQueue()
    Assert.Fail()
}
catch(IOException)
{
}
```

**Test Case 8.2:**

```
Leafqueues.add(name="A1")
Leafqueues.add(name="A2")
try
{
    CapacityScheduler.parseQueue()
}
catch(IOException)
{
    Assert.Fail()
}
```

**2. Verify that the user is authorized to perform an action in the Node Manager.**

**Context** ContainerManagerImpl::authorizeUser()  
**Inv:** remoteUgi.getUserName() = nmTokenIdentifier.getApplicationAttemptId()

**Test Case 9.1:**

```
UserGroupInformation rUgi = new UserGroupInformation()
NMTokenIdentifier nmTokenId = new NMTokenIdentifier()
rUgi.setUserName("RIT")
nmTokenId.setApplicationAttemptId("RIT")
try
{
    CapacityScheduler.parseQueue()
}
catch(RPCUtil.getRemoteException)
{
    Assert.Fail()
}
```

**Test Case 9.2:**

```
UserGroupInformation rUgi = new UserGroupInformation()
```

```

NMTokenIdentifier nmTokenId = new NMTokenIdentifier()
rUgi.setUserName("MIT")
nmTokenId.setApplicationAttemptId("RIT")
try
{
    CapacityScheduler.parseQueue()
    Assert.Fail()
}
catch(RPCUtil.getRemoteException)
{
}

```

## **10.        *Do not support Node Manager recovery with an ephemeral server port.***

**Context** ContainerManagerImpl::serviceStart()

**Inv:** initialAddress.getPort()!=0 or !context.getNMStateStore().canRecover()

### **Test Case 10.1:**

```

initialAddress.setPort(0)
context.getNMStateStore().setCanRecover(True)
try
{
    ContainerManagerImpl.serviceStart()
    Assert.Fail()
}
catch(IllegalArgumentException)
{
}

```

### **Test Case 10.2:**

```

initialAddress.setPort(8180)
context.getNMStateStore().setCanRecover(True)
try
{
    ContainerManagerImpl.serviceStart()
}
catch(IllegalArgumentException)
{
    Assert.Fail()
}

```

### **Test Case 10.3:**

```

initialAddress.setPort(0)
context.getNMStateStore().setCanRecover(False)
try
{
    ContainerManagerImpl.serviceStart()
}

```


```
catch (IllegalArgumentException)
{
    Assert.Fail()
}
```

**Test Case 10.4:**

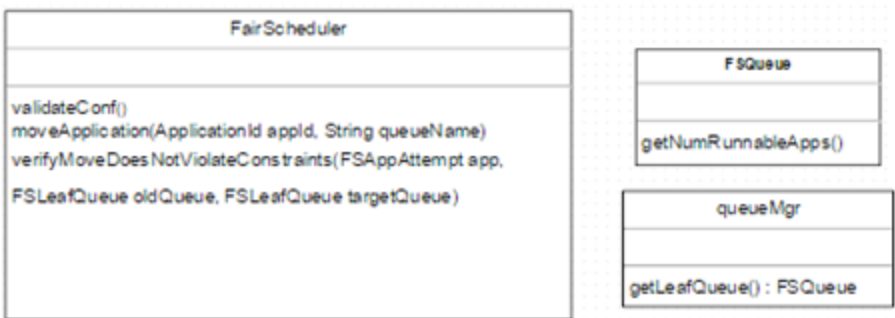
```
initialAddress.setPort(8180)
context.getNMStateStore().setCanRecover(False)
try
{
    ContainerManagerImpl.serviceStart()
}
catch (IllegalArgumentException)
{
    Assert.Fail()
}
```

Design models that can be related to OCL

OCL 1:




OCL 2, 3, 7:




OCL 4, 5:



OCL 6:



OCL 8:



OCL 9, 10:



| ContainerManagerImpl   |
|--|
| initialAddress<br>context  |
| authorizeUser(UserGroupInformation remoteUgi, NMTokenIdentifier nmTokenIdentifier)<br>serviceStart() |

## The Methodology

To generate test cases from OCL we follow the following procedure:

1. Identify the context (and the files involved) of the OCL
2. Find all the possible different combinations of the variables used in the invariance, preconditions or postconditions. For each of these combinations, one test case has to be generated.
3. For each test case, assign values to the variables between the ranges previously selected.
4. Figure out what is the expected value of the function being tested, given the specific circumstances. If it is not the expected result, raise an assert fail exception. In case that the expected result of a function is to raise an exception, this specific exception have to be caught, so the assert exception can behave as desired in the test case.

## Appendix:

### Hadoop YARN architecture recovery process:

For recovering this architecture we followed a top bottom approach.

At the beginning, we were not completely sure if heartbeat was implemented in Hadoop. We thought that it could be a ping echo or a heartbeat and if it was heartbeat we were not sure about the direction in which the heartbeat was sent. To know what actually was happening, we decided to search in the Hadoop issuer tracker and to look up in the code.

We got to the conclusion that they implemented an improved heartbeat in the ResourceTrackerService class. The heartbeat is sent from NM to RM and a response heartbeat message containing the next Interval heartbeat was sent from RM to NM. Even though there are exchanges of messages in both directions, we believe that this should be considered an improved heartbeat instead of a ping echo, because of the intentions. In this case, the Node Manager sends heartbeat to the Resource Manager to let him know that he is still alive and the response message is just to make more efficient the heartbeat by controlling the heartbeat interval. If it were ping echo, the Node Manager would have been using the heartbeat response just to make sure that the Resource manager received its heartbeat. Below we can see links and images that support this hypothesis.

<http://grokbase.com/t/hadoop/yarn-issues/12chd3fmgm/jira-created-yarn-275-make-nodemanager-to-not-blindly-heartbeat-irrespective-of-whether-previous-heartbeat-is-processed-or-not>

Bikas Saha (JIRA) at Dec 28, 2012 at 5:48 pm

[ <https://issues.apache.org/jira/browse/YARN-275?page=com.atlassian.jira.plugin.system.issuetabpanels:comment-tabpanel&focusedCommentId=13540621#comment-13540621> ]

Bikas Saha commented on YARN-275:

I briefly looked at the patch. The general approach seems promising. I have some comments on how we can structure this changes

We could break this work into 2 parts

- 1) protocol changes in heartbeat to transfer heartbeat control frequency from NM to RM. After this, in every heartbeat the RM will tell the NM when to send the next heartbeat. That value can be hardcoded (like it is currently) but preferably we can have an RM config that defines what the minimum heartbeat interval should be and use that. For this part, I don't think we need both backoff and heartbeatinterval in the heartbeat response. We can just have only heartbeatinterval that is always respected by the NM.
- 2) add some logic/heuristic to the RM so that it can dynamically change the heartbeat interval based on its current processing load/rate. This way the interval can be made longer when the RM is not keeping up with heartbeats.

If you think this break-up of works makes sense then we can create 2 sub-tasks under this jira for the 2 parts.

I have some additional ideas on part 1 also.

When a heartbeat comes at time T to the RM then it can choose to

- A) accept the request at time T and ask NM to heartbeat after time T+K with new information. This adds more load to the current RM load. This is what the current code does. So no change is required to do this.
- B) reject the request at time T and ask NM to heartbeat after time T+K with current+new information. This does not increase load on RM but makes NM more complex because it needs to hold onto the last heartbeat data and merge in new data to it.

What do you think about these alternatives?

[Show quoted text](#)

[reply](#) | [permalink](#)

```

346 @Override
347 public NodeHeartbeatResponse nodeHeartbeat(NodeHeartbeatRequest request)
348     throws YarnException, IOException {
349
350     NodeStatus remoteNodeStatus = request.getNodeStatus();
351     /**
352      * Here is the node heartbeat sequence...
353      * 1. Check if it's a registered node
354      * 2. Check if it's a valid (i.e. not excluded) node
355      * 3. Check if it's a 'fresh' heartbeat i.e. not duplicate heartbeat
356      * 4. Send healthStatus to RMNode
357      */
358
359     NodeId nodeId = remoteNodeStatus.getNodeId();
360
361     // 1. Check if it's a registered node
362     RMNode rmNode = this.rmContext.getRMNodes().get(nodeId);
363     if (rmNode == null) {
364         /* node does not exist */
365         String message = "Node not found resyncing " + remoteNodeStatus.getNodeId();
366         LOG.info(message);
367         resync.setDiagnosticsMessage(message);
368         return resync;
369     }
370
371     // Send ping
372     this.rmLivelinessMonitor.receivePing(nodeId);
373
374     // 2. Check if it's a valid (i.e. not excluded) node
375     if (!this.nodesListManager.isValidNode(rmNode.getHostName())) {
376         String message =
377             "Disallowed NodeManager nodeId: " + nodeId + " hostname: "

```